

Project 3 Essay - The Atrophy of Epistemic Agency in a Post-AGI Social Order

Marco Schürer Drews

Democratic societies depend on a capacity that is easy to overlook precisely because it is so ordinary: the ability of citizens to form, revise, and defend their own judgments about how the world works and what ought to be done about it. This capacity -call it epistemic agency- is not a fixed endowment. It is a practiced skill. It develops through the effortful work of evaluating competing claims, tolerating uncertainty long enough to reason through it, and arriving at positions one is prepared to defend in the company of others who have done the same. Landemore (2013) has shown that the epistemic advantage of democratic decision-making depends on two properties of its participants: individual competence and cognitive diversity. Both presuppose that citizens are actually doing the cognitive work. If that presupposition fails, the epistemic case for democracy collapses alongside it.

The argument I want to make is that advanced AI systems threaten epistemic agency - not only limited to producing falsehoods - but by making the practice of independent judgment feel unnecessary. The familiar worry about hallucination and misinformation, while real, understates the problem. A system that is reliably wrong can be caught and corrected. A system that is right often enough to be trusted as a default source of judgment is far more corrosive, because it removes the occasion for the cognitive work through which epistemic agency is sustained. Risko and Gilbert (2016) document this dynamic under the heading of “cognitive offloading”: when external tools reliably handle a cognitive task, the internal capacity to perform that task declines through disuse. Their framework, developed for memory and spatial reasoning, applies with particular force to evaluative judgment. The cost-benefit threshold for independent reasoning shifts; once a fluent and plausible answer is available instantly, the effortful alternative becomes harder to justify to oneself.

Consider what happens when the cost of generating a reasoned position on any question drops to nearly zero. A citizen who might previously have read competing accounts of a policy dispute, weighed the evidence, and formed a tentative view now receives a fluent, well-sourced summary that arrives faster than the first paragraph of their own thinking would have. The output may be

excellent, but the problem is not in the quality, the problem is that the citizen has been converted from a participant in collective reasoning to a consumer of its products. According to Vallor (2015), when technology mediates a skill-dependent activity, the skill itself can atrophy even as the activity's outputs improve. She frames this as "moral deskilling," but the logic extends to any capacity that requires practice, including the capacity to form independent epistemic judgments. Over time, the less capable one becomes of forming independent judgments, the less one notices the absence.

This matters for the social contract in a specific way. Democratic legitimacy rests on the assumption that citizens are capable of the kind of collective deliberation that makes self-governance meaningful. Rawls's (1997) original position presupposes agents who can evaluate principles of justice from behind the veil of ignorance (agents, that is, who possess the capacity for considered judgment). Habermas (1994) grounds communicative legitimacy in the requirement that all affected parties be able to participate in rational discourse, which requires not just the opportunity to speak but the competence to reason. If that competence erodes - if a sufficient proportion of the population delegates its reasoning to AI systems and consequently loses the practiced ability to reason independently - then democratic procedures persist as form while the substance beneath them hollows out. Elections still happen. Legislatures still convene. But the deliberative capacity that makes these institutions more than ritual is diminished.

The mechanism here is convenience rather than plain manipulation. The most plausible path to the erosion of epistemic agency runs through entirely voluntary adoption of tools that make cognitive life easier, no coercion nor deception is required. People will delegate their reasoning for the same reason they delegate their arithmetic, because the machine does it faster and, in most cases, well enough. Lazar (2025) argues that computing technologies have restructured political association in ways analogous to industrialization, creating new forms of intermediary power that challenge existing frameworks of political philosophy. The argument I am building in this essay extends that diagnosis; the restructuring reaches not just the institutional architecture of political life but the cognitive capacities of the citizens who are supposed to govern through it.

Two features of this dynamic make it particularly resistant to governance intervention. i) First, the harm is invisible at the individual level. No single act of consulting an AI assistant constitutes a problem. The damage is cumulative and population-level, visible only in the

declining quality of public deliberation over time. This places it in the category of systemic risks that existing regulatory frameworks are poorly designed to address, since those frameworks are built around identifiable harms to identifiable individuals. And ii) second, the process is self-reinforcing. As epistemic agency declines, the demand for AI-mediated judgment increases, because citizens who have lost confidence in their own reasoning rely more heavily on the systems whose convenience produced the loss. There is no natural equilibrium that arrests the cycle short of a population that has substantially outsourced its capacity for collective sense-making.

If this assessment holds, it implies, at the very least, that any rigorous analysis of social structures in a post-AGI world must confront a stark reality - the individuals participating in the social contract might lack the very cognitive abilities that the contract presupposes. The question is not whether AGI will make people less intelligent - that framing is too crude and too individual. The question is whether a society that has delegated its epistemic labor to artificial systems can sustain the forms of collective agency on which democratic governance depends.

References

- Habermas, J. (1994). Three normative models of democracy. *Constellations*, 1(1), 1–10.
- Landemore, H. (2013). Deliberation, cognitive diversity, and democratic inclusiveness: An epistemic argument for the random selection of representatives. *Synthese*, 190(7), 1209–1231.
- Lazar, S. (2025). Governing the algorithmic city. *Philosophy & Public Affairs*, 53(2), 102–168.
- Rawls, J. (1997). The idea of public reason revisited. *University of Chicago Law Review*, 64(3), 765–807.
- Risko, E. F., & Gilbert, S. J. (2016). Cognitive offloading. *Trends in Cognitive Sciences*, 20(9), 676–688.
- Vallor, S. (2015). Moral deskilling and upskilling in a new machine age: Reflections on the ambiguous future of character. *Philosophy & Technology*, 28(1), 107–124.